

Task 5. Analyzing Neuroimaging Data with PySpark and Thunder

Advances in imaging equipment and automation have led to an overabundance of data on the functioning of the brain. Technologies today can sample brain activity from a large number of neurons in a large region while organisms are actively behaving. For example, by simultaneously recording the electrical activity of every neuron of the mouse brain over an extended period of time, the amount of data generated will create completely new paradigms for biology, that will require the development of tools to extract value from such unprecedented amount of information.

In this Notebook, we use PySpark and the Thunder project (<https://github.com/thunder-project/thunder>), which is developed on top of PySpark, for processing large amounts of time series data in general, and neuroimaging data in particular. We will use these tools for the task of understanding some of the structure of Zebrafish brains, which is a typical (and simple) example used in Neuroimaging. Using Thunder, we will cluster different regions of the brain (representing groups of neurons) to discover patterns of activity as the zebrafish behaves over time.

Note: Please, use the documentation for the Thunder API (<http://docs.thunder-project.org/>) to learn the details of function calls!

Goals

The main goals of this notebook are:

1. Learn about Thunder and how to use it
2. Revisit the K-Means algorithm and the method for choosing K
3. Learn alternative approaches to improve the results

Steps


1. In section 1, we go through some background concepts that are used in this notebook.
2. Next, in section 2, we will get familiar with Thunder, its methods and its data types, by working on some simple tasks.
3. Finally, in section 3, we will build a model to cluster the neurons of a zebrafish based on their behavior. In this step, we will learn about how to use K-Means when the value of K is unknown. Finally, some tricks to improve the results are introduced.

1. Background concepts

In this section, we cover the terminology and the concepts that constitute the domain knowledge for this notebook.

As it should be well-known, a **pixel** is a combination of "picture element": digital images can be modeled as simple 2-dimensional (2D) matrices of intensity values, and each element in the matrix is a pixel. In color images, a pixel contains values of red, green, and blue channels. In a grayscale image, the three channels have the same value, such that each pixel is reduced to be a single value.

A single 2D image is not nearly enough to express 3D objects, which use a **voxel**, representing a value of the 3D image on a regular grid in a three-dimensional space. A possible technique to work on 3D images is to acquire multiple 2D images of different slices (or planes, or layers) of a 3D object, and stack them one on top of each other (a z-stack). This ultimately produces a 3D matrix of intensity values, where each value represents a volume element or voxel.



This z-stack image has 4 layers. A point is a voxel. It can be determined by the layer's index and the position in that layer.

In the context of the Thunder package, we use term `image` for 3D-image or `stack image`. Thunder uses `Image` type to represent 3D-image. Each `Image` object is a collection of either 2D images or 3D volumes. In practice, it wraps an n-dimensional array, and supports either distributed operations via Spark or local operations via numpy, with an identical API.

Stack-images can represent 3D objects, but it can be difficult to take the temporal relationship of the images into account. To do that, we need another data structure that shows the changes of voxels over time. In the Thunder package, the internal `Series` type can be used exactly for this purpose. Each `Series` is a 1D array such that each element is a value of the voxel at a timestamp.

The most common series data is time series data, in which case the index is time and each record is a different signal, like a channel or pixel.

We now have sufficient material to start playing with Thunder !!!

2. Let's play

Well, wait a second before we play... Remember, we're going to use Spark to perform some of the computations related to this Notebook. Now, when you spin a Zoe Notebook application (this comment is valid for students at Eurecom), you'll gain access to an individual, small Spark cluster that is dedicated to your Notebook. This cluster has two worker machines, each with 6 cores. As such, a good idea to obtain smooth performance and a balanced load on the workers, is to repartition your data (i.e., the RDDs you use to represent images or time series).

In this Notebook we **expect** students to take care of repartitioning, and such care will be compensated by bonus points.

2.1. Play with Image objects

a. Loading image data

Both images and series can be loaded from a variety of data types and locations. You need to specify whether data should be loaded in 'local' mode, which is backed by a numpy array, or in 'spark' mode, which is backed by an RDD by using the optional argument `engine`. The argument `engine` can be either `None` for local use or a `SparkContext` for distributed use with Spark.

```
import thunder as td

# Load data from tif images
data = td.images.fromtif('/path/to/tifs')

# Load data from numpy-arrays
data = td.series.fromarray(somearray)
data_distributed = ts.series.fromarray(somearray, engine=sc)
```

We can load some example image data by:

```
In [1]: import thunder as td
import numpy as np

# Load some example image data
image_data = td.images.fromexample('fish', engine=sc)

# print the number of images
print(image_data.count())
```

20

b. Inspecting image data

```
In [2]: %matplotlib inline
import matplotlib.pyplot as plt

# import two function to draw images easier
from showit import image as draw_image
from showit import tile as draw_tile

print("Shape of the data:", image_data.shape)

first_image = image_data.first() # get the values of Image object
# or first_image = image_data[0] # get the Image object

print("Shape of the data of the first image:", first_image.shape)

print("Data of the first image:", first_image)

# draw the first layer of the first image
draw_image(first_image[0])

# draw all layers of the first image
draw_tile(first_image)




# we can use index slices to take images
samples = image_data[0:6]
```

```

Shape of the data: (20, 2, 76, 87)
Shape of the data of the first image: (2, 76, 87)
Data of the first image: [[[26 26 26 ..., 26 26 26]
 [26 26 26 ..., 26 26 26]
 [26 26 26 ..., 27 27 26]
 ...,
 [26 26 26 ..., 27 27 26]
 [26 26 26 ..., 27 26 26]
 [25 25 25 ..., 26 26 26]]]

[[[25 25 25 ..., 26 26 26]
 [25 25 25 ..., 26 26 26]
 [26 26 26 ..., 26 26 26]
 ...,
 [26 26 26 ..., 26 26 26]
 [26 26 26 ..., 26 26 26]
 [25 25 25 ..., 26 26 26]]]

```



From the result above, the shape of the loaded data is (20, 2, 76, 87). It means we have total 20 3D images objects. Each image has 2 layers, each layer has size 76x87.

Note that, although data is not itself an array (it can be a kind of RDD), we can index into it using bracket notation, and pass it as input to plotting methods that expect arrays. In these cases, the data will be automatically converted.

One of the advantages of working in Python is that we can easily visualize our data stored into Spark RDDs using the Matplotlib library. Function `draw_image` and `draw_tile` that take advantages of Matplotlib are examples.

Question 1

- a) Use the function `imshow` from matplotlib to plot each layer of the first image in `image_data`. b) Discuss the choice of parameters you use for the method `imshow`

```
In [3]: img = image_data.first()
# or:
# img = image_data[1]

# show the first layer
plt.imshow(img[0], interpolation='nearest', aspect='equal', cmap='gray')
plt.title("Layer 0 of the first image")
plt.show()

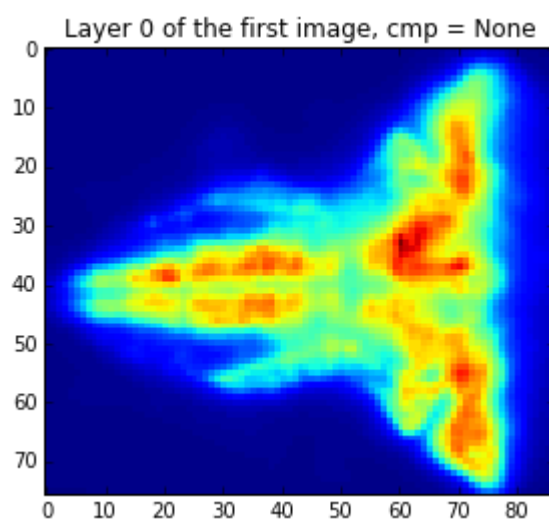

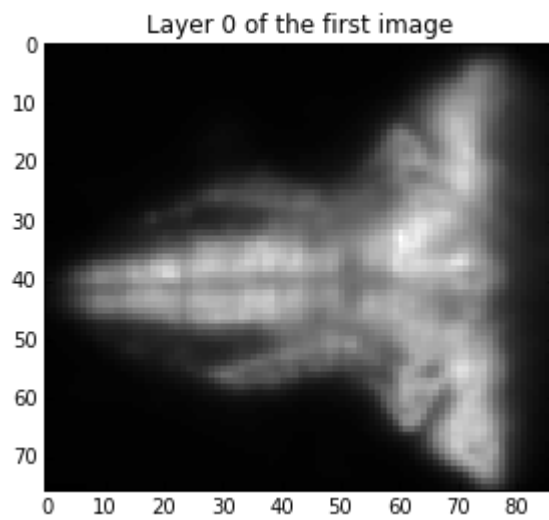
# show the second layer
plt.imshow(img[1], interpolation='nearest', aspect='equal', cmap='gray')
plt.title("Layer 1 of the first image")
plt.show()

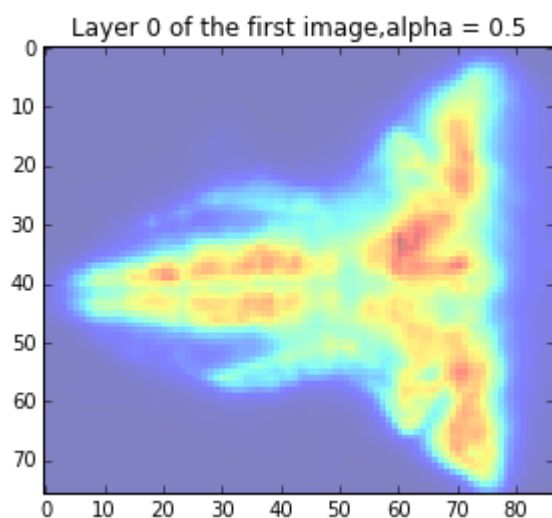
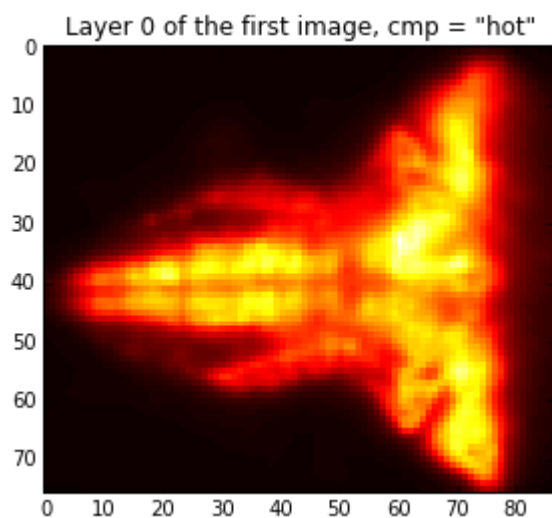
#cmap is set to none
# show the first layer
plt.imshow(img[0], interpolation='nearest', aspect='equal')
plt.title("Layer 0 of the first image, cmp = None")
plt.show()
#cmap os set to "hot"
plt.imshow(img[0], interpolation='nearest', aspect='equal',cmap = "hot")
plt.title("Layer 0 of the first image, cmp = \"hot\"")
plt.show()

#change the blending value alpha
# show the first layer
plt.imshow(img[0], interpolation='nearest', aspect='equal',alpha = 0.5 )
plt.title("Layer 0 of the first image,alpha = 0.5")
plt.show()

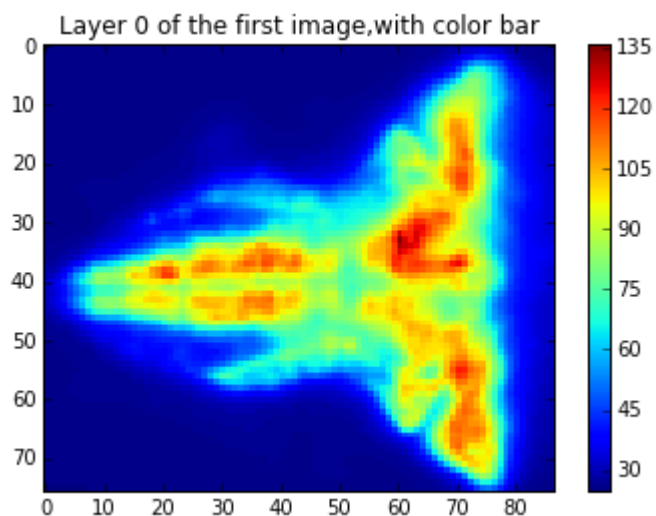
#add the color bar
plt.imshow(img[0], interpolation='nearest', aspect='equal' )
plt.title("Layer 0 of the first image,with color bar")
plt.colorbar()
plt.show()

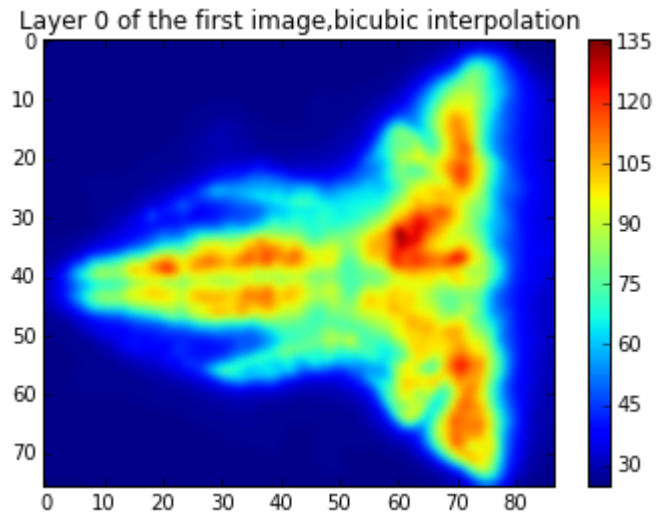
#thumbnail image
plt.imshow(img[0], interpolation='bicubic', aspect='equal' )
plt.title("Layer 0 of the first image,bicubic interpolation")
plt.colorbar()
plt.show()
```





```
/opt/conda/lib/python3.5/site-packages/matplotlib/collections.py:590: FutureWarning: elementwise comparison failed; returning scalar instead, but in the future will perform elementwise comparison
  if self._edgecolors == str('face'):
```





Question 2

Calculate the standard deviation across all images you have in `image_data` (that is, our dataset). To clarify, let's focus on an individual layer (say the first layer of each image). For every `voxel`, compute the standard deviation of its values across different images for the same layer. Visualize the standard deviation you obtain, for example concerning a single layer (as before, say the first layer).



HINT 1 to avoid wasting time and energy, make sure you lookup for methods that could help answer the question from the Thunder documentation.

HINT 2 We can also use function `draw_image(<data>)` to plot an image in a simple way instead of using many statements with matplotlib as before.

NOTE Comment the image you obtain. What does it mean to display the standard deviation across all images in a single layer?

```
In [4]: #####@SOLUTION@!#####  
# calculate standard deviation of images  
std_imgs = image_data.first()  
image_std = image_data.std()  
# standard deviation of Layer 0 of the first image  
image_std_0 = image_std.first()  
draw_image(image_std_0[0])  
# standard deviation of Layer 1 of the first image  
image_std_1 = image_std.first()  
draw_image(image_std_1[1])
```

```
Out[4]: <matplotlib.image.AxesImage at 0x7fb9a0127b00>
```



In the figures shown above, we can see the different distribution of the standard deviations across the image. The stddev stands for the activation state of a single neuron. The higher the stddev is, the more likely that this neuron is activated. On the contrary, a lower value of shows that this neuron is more likely in a unactivated state.

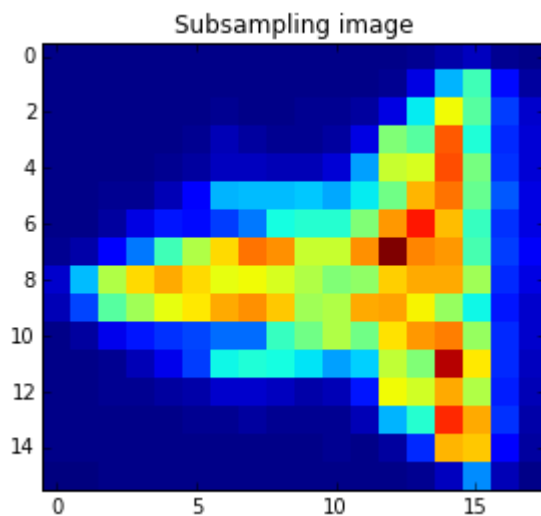
c. Selecting samples of image data

The Images API offers useful methods for working with large image data. For example, in some cases it is necessary to subsample each image, to make sure we can analyze it efficiently.

Question 3

The source code below subsamples image data with different ratios on different dimensions. a) Complete the source code to plot the first layer of the first image. b) What is the shape of `image_data` before and after subsampling?

```
In [5]: subsampled = image_data.subsample((1, 5, 5))
# Stride to use in subsampling. If a single int is passed, each dimension of the image
# will be downsampled by this same factor. If a tuple is passed, it must have the same
# dimensionality of the image. The strides given in a passed tuple will be applied to
# each image dimension
plt.imshow(subsampled.first()[0], interpolation='nearest', aspect='equal', cmap=None)
plt.title("Subsampling image")
plt.show()
print("Before subsampling:", image_data.shape)
print("After subsampling:", subsampled.shape)
```



Before subsampling: (20, 2, 76, 87)
 After subsampling: (20, 2, 16, 18)

Note that `subsample` is an RDD operation, so it returns immediately. Indeed, we know that in Spark you must apply a RDD action to trigger the actual computation.

d. Converting image data

We can also convert an RDD of images to a RDD of series by:

```
In [6]: seriesRDD = image_data.toseries()  
seriesRDD.cache()
```

```
Out[6]: Series  
mode: spark  
dtype: uint8  
shape: (2, 76, 87, 20)
```

Question 4

According to your understanding about `Series` objects which was introduced in section 1, what is the shape of `seriesRDD` and its elements ? Comment your results, don't just display numbers.

```
In [7]: seriesRDD.shape
```

```
Out[7]: (2, 76, 87, 20)
```

The shape of `seriesRDD` is `(2,76,87,20)`, which means we have 20 time nodes. The data point in each time node has the shape of `(2,76,87)`, indicating that the image has `2 layers`, and each layer consists of a `76*87` pixels(voxels).

For a large data set that will be analyzed repeatedly as a `Series`, it will ultimately be faster and more convenient to save Images data to a collection of flat binary files on a distributed file system, which can in turn be read back in directly as a `Series`, rather than repeatedly converting the images to a `Series` object. This can be performed either through a `ThunderContext` method, `convertImagesToSeries`, or directly on an `Images` object, as done below:

We will study about `Series` object in the next section.

2.2. Play with Series objects

a. Loading Series data

In this section, we use a sample data to explore Series objects.

```
In [8]: # series_data = td.series.fromexample('fish', engine=sc)
# series_data = td.series.frombinary(path='s3n://thunder-sample-data/series/fish', engine=sc)
series_data = image_data.toseries()
```

b. Inspecting Series data

Series_data is a distributed collection of key-value records, each containing a coordinate identifier and the time series of a single voxel. We can look at the first record by using `first()`. It's a key-value pair, where the key is a tuple of int (representing a spatial coordinate within the imaging volume) and the value is an one-dimensional array.

```
In [9]: first_series = series_data.first() # get the values of Series object
#first_series = series_data[0] # get a Series object

print("Shape of series:", series_data.shape)
print("The first series:", first_series)
print("Each element in series has", len(first_series), "values")

# print the 10th value of voxel (0,0,0)
# layer = 0
# coordinator = (0,0) in that layer
print("value 10th of voxel (0,0,0):", np.array(series_data[0,0,0,10]))

Shape of series: (2, 76, 87, 20)
The first series: [26 26 26 26 26 26 26 25 26 25 25 25 26 26 26 26 26 26 2
6]
Each element in series has 20 values
value 10th of voxel (0,0,0): 25
```

The loaded series data is a multi-dimensional array. We can access the values of a voxel in time series by using a tuple as above. In our data, each voxel has 20 values corresponding to 20 states at 20 different times.

c. Selecting Series data

Series objects have a 1D index, which can be used to subselect values.

```
In [10]: print("shape of index:", series_data.index.shape)
print("the first element of a subset", series_data.between(0,8).first())

shape of index: (20,)
the first element of a subset [26 26 26 26 26 26 26 25]
```

Values can be selected based on their index:

```
In [11]: print(series_data.select(lambda x: x > 3 and x < 8).index)
print(series_data.select(lambda x: x > 3 and x < 8).first())


[4, 5, 6, 7]
[26 26 26 25]
```

Question 5

Plot the first 20 values of **all** series objects (that is the values of a voxel) in the series data. This means, on the same plot, you should visualize the values each voxel takes in the first 20 time intervals.

```
In [12]: import numpy as np

# only select the first 20 states of each object
samples = series_data.between(0,20).tordd().values().collect()
plt.figure(figsize=(20,10))
plt.plot(np.array(samples).T)
plt.xlabel("picture number")
plt.ylabel("picture value")
plt.show()
```



Now, another objective we can have is to select specific series objects within the same series data. For example, we can select objects randomly by using function `sample`.

Question 6

Let's plot a random subset of the data using the method `sample`. Complete the source code below to plot the first 20 values of 30 objects that are selected randomly among those that pass the condition on the standard deviation, using function `sample`.


```
In [13]: print(series_data.map(lambda x : x.std()).max().first())
series_data.shape
```

```
[ 20.80931282]
```

```
Out[13]: (2, 76, 87, 20)
```

```
In [14]: # select 30 objects randomly which have standard deviation > threshold
# Extract random subset of records, filtering on a summary statistic.
examples = series_data.filter(lambda x: x.std() > 1.0).sample(30)
# only plot first 20 states of each object
plt.figure(figsize=(20,10))
plt.xlabel("picture number",fontsize = 20)
plt.ylabel("picture value",fontsize = 20)
plt.plot(np.array(examples).T,linewidth = 2.0)

plt.show()
```




d. Preprocessing Series data


A Series objects has some methods which can be useful in an eventual preprocessing phase.

For example, `center` subtracts the mean, `normalize` subtracts and divides by a baseline (either the mean, or a percentile).

```
In [15]: #sustract the mean
examples = series_data.center().filter(lambda x: x.std() >= 10).sample(50)
plt.figure(figsize=(20,10))
plt.xlabel("picture number",fontsize = 20)
plt.ylabel("picture value",fontsize = 20)
plt.plot(np.array(examples).T)
plt.show()
```




```
In [16]: normalizedRDD = series_data.normalize(method='mean').filter(lambda x: x.std()
>= 0.1).sample(50)
plt.figure(figsize=(20,10))
plt.xlabel("picture number",fontsize = 20)
plt.ylabel("mean",fontsize = 20)
plt.plot(np.array(normalizedRDD).T)
plt.show()
```



e. Computing statistics about Series data

A Series can be summarized with statistics both within and across images. To summarize **across records** (the statistic of all voxels at each timestamp), we can do the following:

```
In [17]: plt.figure(figsize=(15,10))
plt.xlabel("picture number",fontsize = 20)
plt.ylabel("statistics value",fontsize = 20)
max_plot, = plt.plot(series_data.normalize().max(),label = "max",linewidth =
2.0);
mean_plot, = plt.plot(series_data.normalize().mean(),label = "mean",linewidth
= 2.0);
min_plot, = plt.plot(series_data.normalize().min(),label = "min",linewidth =
2.0);
plt.legend(handles = [max_plot,mean_plot,min_plot] )
plt.show()
```



To summarize **within records**, we can use the map method:

```
In [18]: means = series_data.map(lambda x: x.mean())
flat_means = means.flatten().toarray()
flat_stdevs = stdevs = series_data.map(lambda x: x.std()).flatten().toarray()
print("means:", flat_means)
print("length of means:", len(flat_means))
print("mean of the first series:", flat_means[0])
print("standard deviation of the first series:", flat_stdevs[0])

means: [ 25.8  25.85 25.7  ..., 26.   26.   26. ]
length of means: 13224
mean of the first series: 25.8
standard deviation of the first series: 0.4
```

means is now a Series object, where the value of each record is the mean across the time series for that voxel.


Note that in the source code above, we use function `toarray` to return all records to the driver as a numpy array.

For this Series, since the keys correspond to spatial coordinates, we can pack the results back into a local array in **driver node**.

To look at this array as an image, we can use function `draw_image` as before.

```
In [19]: # we should recover the shape of means before plotting  
# draw the standard deviations of series that belong to the first layer  
draw_image(flat_means.reshape((2, 76, 87)) [0,:,:])
```

```
Out[19]: <matplotlib.image.AxesImage at 0x7fb9833067f0>
```



Note that `toarray` is an example of a local operation, meaning that all the data involved will be sent to the Spark driver node. In this case, packing the mean is no problem because its size is quite small. But for larger data sets, this can be **very problematic**. So, it's a good idea to downsample, subselect, or otherwise reduce the size of your data before attempting to pack large image data sets!

f. Identifying correlations

In several problem domains, it may also be beneficial to assess the similarity between a designated signal (time series) and another signal of interest by measuring their correlation. For example, say we have two time series corresponding to the consumption of Coca Cola and Pepsi, it would perhaps be interesting to verify whether behavioural patterns are similar for both brands over time.

Simply as a proof of concept, we shall compare our data to a random signal and we expect that, for a random signal, the correlation should be low. The signal can be stored as a numpy array or a MAT file containing the signal as a variable. Note that the size of the signal must be equal to the size of each Series element.

```
In [22]: from numpy import random
         signal = random.randn(len(first_series))
         print("The correlation of the first element with random signal:" ,
               series_data.correlate(signal).first())

         first_element = series_data.first()
         corr = series_data.correlate(np.array(first_element)).first()
         print("The correlation of the first element with itself:", corr)
```

```
The correlation of the first element with random signal: [-0.05607361]
The correlation of the first element with itself: [ 1.]
```

3. Usecase

3.1. Context

Neurons have a variety of biochemical and anatomical properties. Classification methods are thus needed for clarification of the functions of neural circuits as well as for regenerative medicine. In this usecase, we want to categorize the neurons in a fish brain, based on their behavior. The behavior of a neuron can be expressed by the change of its states. The activities of the brains are captured over time into images.

Neurons have a variety of biochemical and anatomical properties. Classification methods are thus needed for clarification of the functions of neural circuits as well as for regenerative medicine. In this usecase, we want to categorize the neurons in a fish brain, based on their behavior. The behavior of a neuron can be expressed by the change of its states. The activities of the brains are captured over time into images.

In this notebook, we use K-Means, a well known clustering algorithm which is also familiar to you, as it was introduced during the last lecture on anomaly detection.

3.2 Data

The dataset we will use is the time series data which we played with in the previous section. Refer to section 2 if you want to duplicate the code to load such data.

3.3. Building model

a. Importing required modules

```
In [23]: %matplotlib inline
import matplotlib.pyplot as plt
from pyspark.mllib.clustering import KMeans, KMeansModel
from matplotlib.colors import ListedColormap
```

b. Loading & inspecting the data

Question 7

Load example series data from `fish`, then normalize and cache it to speed up repeated queries. Print the dimensional information of the loaded data.

```
In [24]: #####@SOLUTION@!#####
# we must normalize it to get best clustering
# image_data = td.images.fromexample('fish', engine=sc)
# data = image_data.toseries()
import thunder as td
data = td.images.fromexample('fish',engine=sc).toseries().normalize()

# cache it to speed up related queries
data.cache()

# check the dimensions oins, patches = plt.hist(stddevs,bins = 19)

print (data.shape)

(2, 76, 87, 20)
```

Question 8

When studying the properties of large data set, we often take a small fraction of it. We have many strategies to select this subset, such as selecting randomly, selecting elements that has the standard deviation bigger than a threshold, or mixing the conditions. In this notebook, we will use the second method as a small demonstration.

In order to choose a good value for the threshold of standard deviation, we should compute the stddev of each series and plot a histogram of a 10% sample of the values.

Complete the source code below to compute the standard deviation of series in data. Plot the histogram of it and discuss it in details. In your opinion, what should be the best value for the threshold ?


```


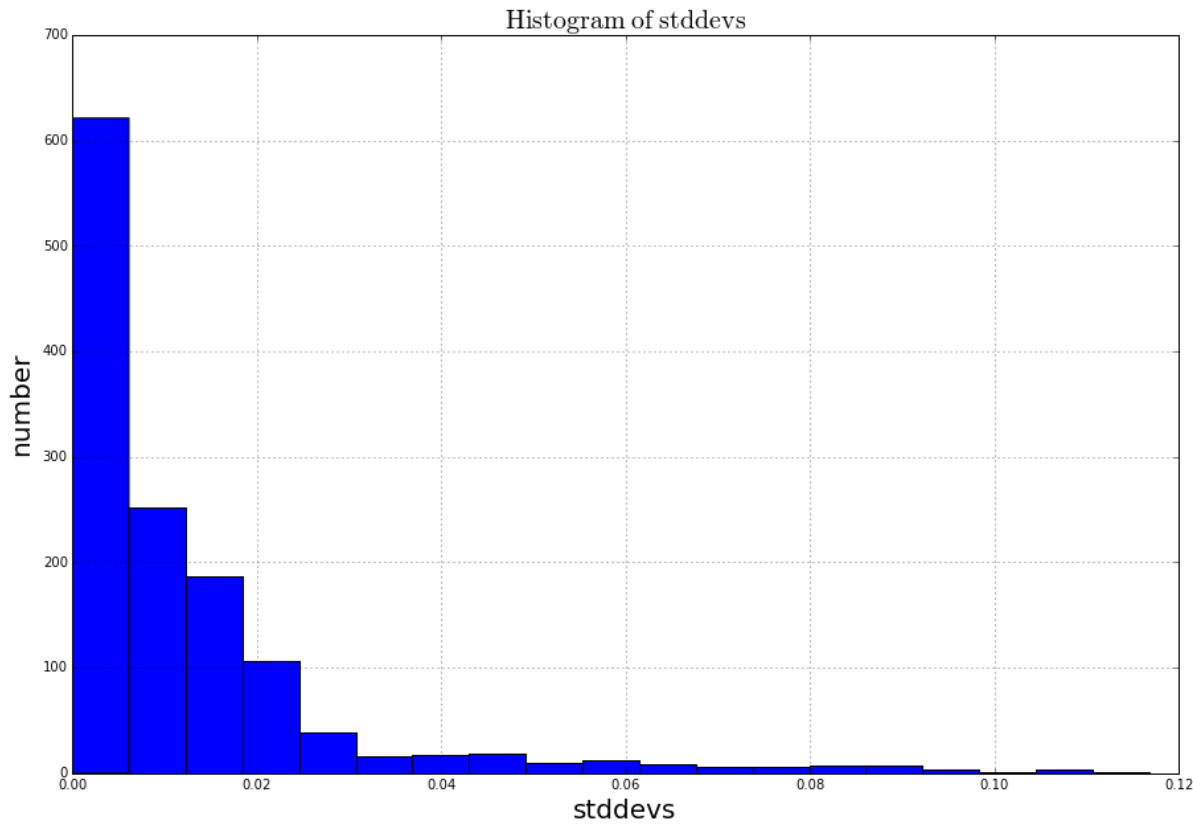
In [114]: # calculate the standard deviation of each series
# then select randomly 10% of value to plot the histogram
from scipy.stats import gaussian_kde
n = data.count()
stddevs = (data.map(lambda x : x.std()).sample(int(n/10)))

# plot the histogram of 20 bins
plt.figure(figsize=(15,10))
h1 = plt.hist(stddevs.flatten().toarray(),bins = 19)

plt.xlabel('stddevs',fontsize = 20)
plt.ylabel('number',fontsize = 20)
plt.title(r'$\mathrm{Histogram\ of\ stddevs}\ $',fontsize = 20)
plt.grid(True)
#plt.plot(stddevs)
#plt.show()
density=gaussian_kde(stddevs.flatten().toarray())
#rint (stddevs)
#print(stddevs.flatten().toarray())
def threshold(percentage,frequence):
    for i in range(0,len(frequence)):
        if sum(frequence[0:i])>=percentage*sum(frequence):
            return i
fig = plt.figure(figsize = (15,10))
plt.plot(np.linspace(0,0.12,100),density(np.linspace(0,0.12,100)),color="black")
print (sum(density(np.linspace(0,0.12,100))))
k= threshold(0.95,h1[0])
print (h1[1][k-1])

```

657.503502226
0.0430312721059



In order to have a better visual performance, here we use the ``kde function`` as the supplement to the histogram

In order to set a rule to calculate the threshold of std, we adopt the following method: 1.Choose the part considered(observed as dark voxel according to human eyes) to be the background of the image from the whole data 2.Calculate the average std of them among the 20 time nodes; 3.Use this average value as the threshold of std.

1. Choose the background area

By plotting the dark area, we choose the [0:19] * [0:19] and [62:75] * [0:17] parts of one image, for the reason that these areas in first image appear to be dark(background). The code and the plot result are shown as below.


```
In [26]: draw_image(first_image[0,0:19,0:19])
```

```
Out[26]: <matplotlib.image.AxesImage at 0x7fb981a1a4a8>
```



```
In [27]: draw_image(first_image[0,62:75,0:17])
```

```
Out[27]: <matplotlib.image.AxesImage at 0x7fb980c540f0>
```



2. Calculate the average of the standard deviation of the chosen area.

```
In [28]: stddevs = data.map(lambda x : x.std()).toarray().reshape(2, 76, 87)
stddevs_background0_a = stddevs[0,0:19,0:19]
stddevs_background0_b = stddevs[0,62:75,0:17]
stddevs_background1_a = stddevs[1,0:19,0:19]
stddevs_background1_b = stddevs[1,0:19,0:19]
stddevs_average1 = np.mean(stddevs_background0_a)
stddevs_average2 = np.mean(stddevs_background0_b)
stddevs_average3 = np.mean(stddevs_background1_a)
stddevs_average4 = np.mean(stddevs_background1_b)
stddevs_average =
np.mean([stddevs_average1,stddevs_average2,stddevs_average3,stddevs_average4])
print("average stddevs is:",stddevs_average)
```

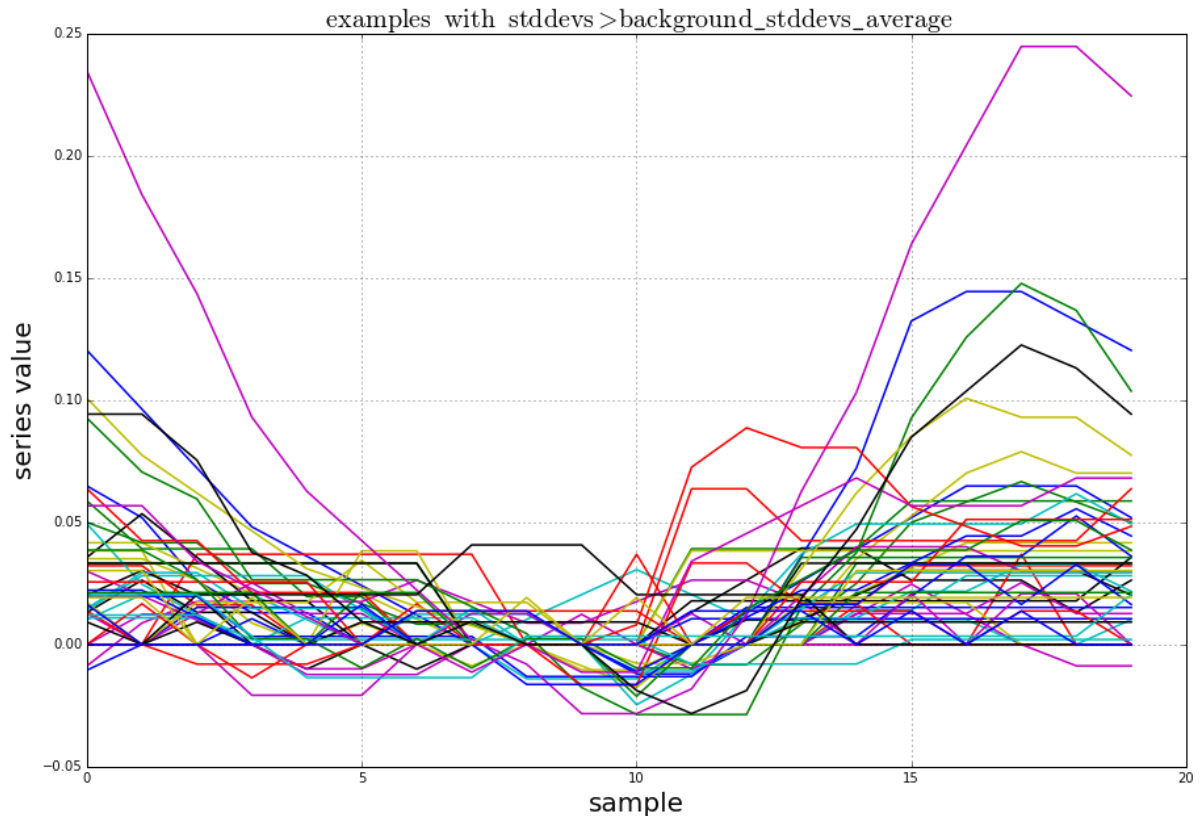
```
average stddevs is: 0.000304580735977
```

Question 9

Extract some samples just to look at the typical structure of the time series data. The objects are selected randomly, and has the standard deviation bigger than the threshold which you picked in question 8. Plot the samples and discuss your obtained figure.

```
In [29]: # sample 50 objects of the data randomly base on the standard deviation
examples = data.filter(lambda x : x.std()>stddevs_average).sample(50)

# plot the sample data
plt.figure(figsize=(15,10))
plt.plot(np.array(examples).T,linewidth = 1.5)
plt.xlabel('sample',fontsize = 20)
plt.ylabel('series value',fontsize = 20)
plt.title(r'$\mathrm{\ examples\ \ with\ \ stddevs > background\_stddevs\_average}$',fontsize = 20)
plt.grid(True)
plt.show()
```



c. Clustering series

In this section, we will use K-means to cluster the series. In other words, we cluster the voxels based on their behavior. Currently, we have no clue about how many groups K of neural behavior. To this end, instead of choosing a single value K , we use multiple values, build model with each K and compare the resulting error values. After that, we can choose the best value of K .

Question 10

Complete the source below to build multiple models corresponding to multiple values of `K` using algorithm KMeans of Thunder. a) Comment the structure of the code. Precisely, focus on the `for` loop, and state what is parallel and what is not. b) Can you modify the structure of the code such that you use the most of the parallelization capabilities of Spark?

```
In [30]: from sklearn.cluster import KMeans
import numpy as np
# declare the possible values of K
ks = [5, 10, 15, 20, 30, 50, 100, 200]

# convert series data to rdd of values
training_data = data.tordd().map(lambda x: np.array(x[1])).cache()

def buildModels(data):
    # declare the collection of models
    models = []
    # build model for each K and append to models
    for k in ks:
        models.append(KMeans(n_clusters=k).fit(data.reshape(data.count(),20)))

        #models.append(KMeans.train(data, k, maxIterations=10, initializationMo
de="random"))
    return models

models = buildModels(data)

/opt/conda/lib/python3.5/site-packages/sklearn/utils/fixes.py:64: Deprecation
Warning: inspect.getargspec() is deprecated, use inspect.signature() instead
if 'order' in inspect.getargspec(np.copy)[0]:
```

d. Testing models & choosing the best one

Next, we evaluate the quality of each model. We use two different error metrics on each of the clusterings.

- The first is the sum across all time series of the Euclidean distance from the time series to their cluster centroids.
- The second is a built-in metric of the KMeansModel object.

Question 11

a) Write function `model_error_1` to calculate the sum of Squared Euclidean Distance from the Series objects to their clusters centroids. b) Comment the choice of the error function we use here. Is it a good error definition?

```
In [31]: from scipy.spatial.distance import cdist
from operator import add

# calculate the Euclidean distance
# from each time series to its cluster center
# and sum all distances
def model_error_1(data, model):
    #result = data.map(lambda x : cdist(np.array([x]),np.array([model.centers
    [model.predict(x)]]),'euclidean')).reduce(add)

    result = data.map(lambda x : cdist(np.array([x]),np.array(model.cluster_centers_
    [model.predict(x)]), 'euclidean')).reduce(add)
    return result[0]

errors = []
for model in models:
    errors.append(model_error_1(training_data,model))
```

```
In [32]: errors
```

```
Out[32]: [array([ 532.18770541]),
array([ 424.96079487]),
array([ 390.84857575]),
array([ 367.47169278]),
array([ 336.88764199]),
array([ 308.23861245]),
array([ 275.52417271]),
array([ 242.21174666])]
```

The main drawback of the error function is that the outlet point will cause a severe impact on the error, so that the result will be extremely huge due to some of these outlet points.

Question 12

a) Write function `model_error_2` to calculate the total of similarity of `Series` objects based on how well they match the cluster they belong to, and then calculate the error by inverse the total similarity. b) Similarly to the previous question, comment the choice of the similarity function.

```

In [32]: # calculate the total of similarity of the model on timeseries objects
# and calculate the error by inverse the total similarity
# Estimate similarity between a data point and the cluster it belongs to.
def similarity(centers, p):
    if np.std(p) == 0:
        return 0
    return np.corrcoef(centers[np.argmin(cdist(centers, np.array([p])))], p)
    [0, 1])

def model_error_2(data, model):
    #similarity_value = data.map(lambda x : similarity(model.centers[model.pre
dict(x)],x)).reduce(add)
    similarity_value = data.map(lambda x : similarity(model.cluster_centers_[m
odel.predict(x)],x)).reduce(add)
    return 1/similarity_value

errors2 = []
for model in models:
    errors2.append(model_error_2(training_data, model))
errors2

```

```

Out[32]: [0.00022349145880541982,
0.00019419607173261889,
0.00019010733737829505,
0.00018341018848375219,
0.00017280730252757808,
0.00016793742771217181,
0.00015773151445360581,
0.00015121480863341662]

```

By inverting the similarity as the error, the error is normalized to [0,1], so that we can avoid the problem of error explode due to the outlier points.

Question 13


Plot the error of the models along with the different values of K in term of different error metrics above. From the figure, in your opinion, what is the best value for 'K' ? Why ?


```
In [33]: def testAndPlotTheResult(data, models):
# compute the error metrics for the different resulting clusterings
# errors of models when using function Sum Square Distance Error
errors_1 = np.asarray([model_error_1(data, model) for model in models])

# error of models when using similarity
errors_2 = np.asarray([model_error_2(data, model) for model in models])

# plot the errors with each value of K
plt.figure(figsize=(20,13))
plt.plot(
    ks, errors_1 / errors_1.sum(), 'k-o',
    ks, errors_2 / errors_2.sum(), 'b:v')
plt.xlabel("k",fontsize=20)
plt.ylabel("error",fontsize=20)
plt.show()

testAndPlotTheResult(training_data, models)
```



According to the figure above, we choose $k = 30$ as the best model. We can see that the both error curve decrease sharply at first then goes to flat. We want to choose the k that have a low error but we don't expect a high k value, which will result in too many clusters and get a overfitting result. Based on this idea, we choose $k = 30$ because it sits between the sharp curve and the flat curve, which makes a tradeoff of the accuracy and overfitting.


Determining the optimal k is particularly troublesome for the k -Means algorithm because error measures based on distance decrease monotonically as k increases. This arises because when k is increased, each cluster is decomposed into more and more clusters, such that each point becomes closer to its cluster mean. In fact, in the extreme case where $k = N$, each point will be assigned to its own cluster, and all distances are reduced to nil. Cross-validation or using holdout data is also unlikely to be particularly effective in this case.

To this end, it is often worth assessing a model by measuring its impact on the overall aim of carrying out the clustering. For example, if we are carrying out k -means for grouping customers having similar taste and purchase history with the ultimate intent of making recommendations to customers, our objective function should measure how effective the recommendations are (perhaps using holdout data). An appealing aspect of using such a metric is that it is no longer guaranteed to behave monotonically with respect to k . We shall investigate this further in Question 20.

Question 14

Plot the centroids of the best model. Do you think that the result is good ?

```
In [34]: # plot the best performing model
bestModel = models[4]
plt.figure(figsize = (15,10))
plt.plot(bestModel.cluster_centers_.T)
plt.xlabel("sequence",fontsize = 20)
plt.ylabel("value",fontsize = 20)
plt.show()
```



As is shown in the figure above, we plot the centroids along the x axis, but we get few information from this figure. In order to have a better visualization, we use PCA method to project the 20 dimensions to 3 and 2 dimensions, to see the positions of each clusters.

Plot the centroids using PCA dim = 3

```
In [35]: from pyspark.mllib.feature import PCA as PCAmlib
from pyspark.mllib.linalg import Vectors
rdd = sc.parallelize([
    Vectors.dense(item) for item in bestModel.cluster_centers_]
    #Vectors.dense([2, 0, 1]),
    #Vectors.dense([0, 1, 0])])

model = PCAmlib(3).fit(rdd)
transformed = model.transform(rdd)

centroids=transformed.collect()
xs=[]
ys=[]
zs=[]
for item in centroids:
    xs.append(item[0])
    ys.append(item[1])
    zs.append(item[2])
```


```
In [36]: from mpl_toolkits.mplot3d import Axes3D

fig = plt.figure(figsize = (15,10))
ax = fig.add_subplot(111, projection='3d')

ax.scatter(xs,ys,zs)
```

/opt/conda/lib/python3.5/site-packages/matplotlib/collections.py:590: FutureWarning: elementwise comparison failed; returning scalar instead, but in the future will perform elementwise comparison
 if self._edgecolors == str('face'):

Out[36]: <mpl_toolkits.mplot3d.art3d.Path3DCollection at 0x7fb9240d8080>



Plot the centroids using PCA dim = 2

```

In [37]: from pyspark.mllib.feature import PCA as PCAmlib
from pyspark.mllib.linalg import Vectors
rdd = sc.parallelize([
    Vectors.dense(item) for item in bestModel.cluster_centers_])
    #Vectors.dense([2, 0, 1]),
    #Vectors.dense([0, 1, 0])])

model = PCAmlib(2).fit(rdd)
transformed = model.transform(rdd)

centroids=transformed.collect()
xs=[]
ys=[]
#zs=[]
for item in centroids:
    xs.append(item[0])
    ys.append(item[1])
#    zs.append(item[2])
from mpl_toolkits.mplot3d import Axes3D

fig = plt.figure(figsize = (15,10))
ax = fig.add_subplot(111)

ax.scatter(xs,ys)

```


Out[37]: <matplotlib.collections.PathCollection at 0x7fb8f413b6d8>

/opt/conda/lib/python3.5/site-packages/matplotlib/collections.py:590: FutureWarning: elementwise comparison failed; returning scalar instead, but in the future will perform elementwise comparison

```

if self._edgecolors == str('face'):

```



e. Visualizing the result

We can also plot an image of labels of neurons, such that we can visualize the group of each neuron.


Question 15

Complete the source code below to visualize the result of clustering.

```
In [38]: # predict the nearest cluster id for each voxel in Series
data_flattened = data.reshape(data.count(),20)
labels = bestModel.predict(data_flattened)

# collect data to the driver
imgLabels = labels.reshape(2,76,87)
# consider the voxel of the first layers
draw_image(imgLabels[0,:,:])
```

Out[38]: <matplotlib.image.AxesImage at 0x7fb8f40b7208>



With the default color scheme, this figure is quite difficult to understand and to distinguish the groups according to their similar colors. So, we should have a smarter color selection. The fact is, when we do clustering, it is often the case that some centers are more similar to one another, and it can be easier to interpret the results if the colors are chosen based on these relative similarities. The method `optimize` tries to find a set of colors such that similarities among colors match similarities among an input array (in this case, the cluster centers). The optimization is non-unique, so you can run multiple times to generate different color schemes.


```

In [39]: from numpy import arctan2, sqrt, pi, abs, dstack, clip, transpose, inf, \
        random, zeros, ones, asarray, corrcoef, allclose, maximum, add, multiply, \
        nan_to_num, copy, ndarray, around, ceil, rollaxis

# these functions below are inspired mainly from Thunder-Project source code,
# v.0.6
# url: https://raw.githubusercontent.com/thunder-project/thunder/branch-0.6/thunder/viz/colorize.py

# Optimal colors based on array data similarity.
def optimize_color(mat):
    mat = np.asarray(mat)

    if mat.ndim < 2:
        raise Exception('Input array must be two-dimensional')

    nclrs = mat.shape[0]

    from scipy.spatial.distance import pdist, squareform
    from scipy.optimize import minimize

    distMat = squareform(pdist(mat, metric='cosine')).flatten()

    optFunc = lambda x: 1 - np.corrcoef(distMat, squareform(pdist(x.reshape(nclrs, 3), 'cosine')).flatten()))[0, 1]
    init = random.rand(nclrs*3)
    bounds = [(0, 1) for _ in range(0, nclrs * 3)]
    res = minimize(optFunc, init, bounds=bounds, method='L-BFGS-B')
    newClrs = res.x.reshape(nclrs, 3).tolist()

    from matplotlib.colors import ListedColormap

    newClrs = ListedColormap(newClrs, name='from_list')

    return newClrs

# Blend two images together using the specified operator.
def blend(img, mask, op=add):
    if mask.ndim == 3:
        for i in range(0, 3):
            img[:, :, i] = op(img[:, :, i], mask)
    else:
        for i in range(0, 3):
            img[:, :, i] = op(img[:, :, i], mask)
    return img

def _prepareMask(mask):
    mask = asarray(mask)
    mask = clip(mask, 0, inf)

    return mask / mask.max()

# Colorize numerical image data.
def transform(cmap, img, mask=None, mixing=1.0):
    from matplotlib.cm import get_cmap
    from matplotlib.colors import ListedColormap, LinearSegmentedColormap,

```



```

hsv_to_rgb, Normalize

img = asarray(img)
dims = img.shape

if cmap not in ['polar', 'angle']:

    if cmap in ['rgb', 'hv', 'hsv', 'indexed']:
        img = copy(img)
        for i, im in enumerate(img):
            norm = Normalize(vmin=None, vmax=None, clip=True)
            img[i] = norm(im)

    if isinstance(cmap, ListedColormap) or isinstance(cmap, str):
        norm = Normalize(vmin=None, vmax=None, clip=True)
        img = norm(copy(img))

if mask is not None:
    mask = _prepareMask(mask)

if isinstance(cmap, ListedColormap):
    if img.ndim == 3:
        out = cmap(img)
        out = out[:, :, :, 0:3]
    if img.ndim == 2:
        out = cmap(img)
        out = out[:, :, 0:3]
else:
    raise Exception('Colorization method not understood')

out = clip(out, 0, 1)

if mask is not None:
    out = blend(out, mask, multiply)

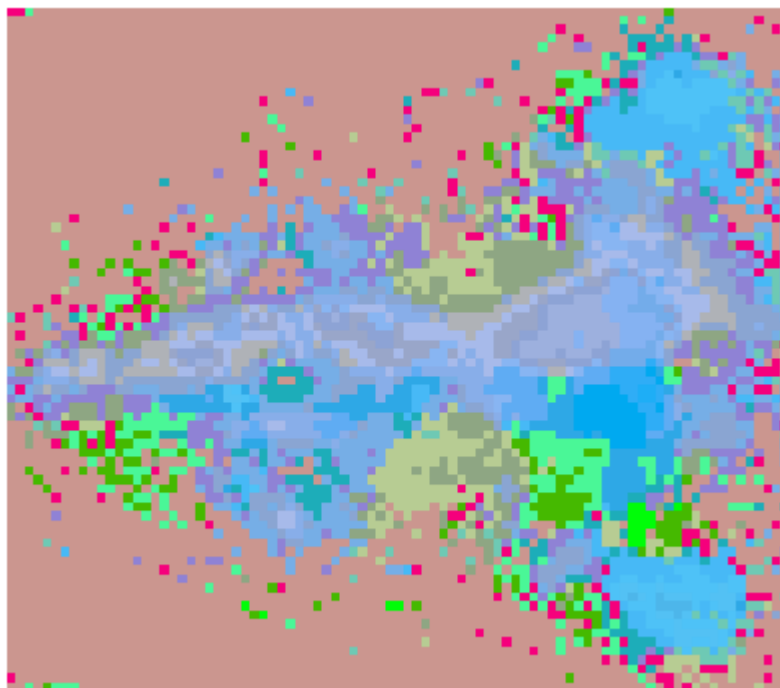

return clip(out, 0, 1)

# generate the better color scheme
newClrs = optimize_color(bestModel.cluster_centers_)
plt.gca().set_color_cycle(newClrs.colors)
plt.plot(np.array(bestModel.cluster_centers_).T);

# draw image with the new color scheme
brainmap = transform(newClrs, imgLabels[0,:,:])
draw_image(brainmap)


```

Out[39]: <matplotlib.image.AxesImage at 0x7fb8e057a978>



```
In [40]: draw_image( transform(newClrs, imgLabels[1,:,:]))
```

```
Out[40]: <matplotlib.image.AxesImage at 0x7fb8e0482f98>
```



f. Improving the result by removing noise

One problem with what we've done so far is that clustering was performed on all time-series without data pre-processing. Many of time-series objects were purely noise (e.g. those outside the brain), and some of the resulting clusters capture these noise signals. A simple trick is to perform clustering after subselecting pixels based on the standard deviation of their time series. First, let's look at a map of the standard deviation, to find a reasonable threshold that preserves most of the relevant signal, but ignores the noise.


Question 16

Try with different threshold of standard deviation to filter the noise. What is the "best value" that preserves most of the relevant signal, but ignores the noise ? Why ?

```
In [41]: # calculate the standard deviation of each voxel  
# then collect to the driver  
stdMap = data.map(lambda x:x.std()).toarray()  
  
# here we should try with many different values of threshold  
# and choosing the best one  
# visualize the map of the standard deviation after filtering  
for thresh in [0.001,0.003,0.007,0.01,0.015,0.02]:  
    draw_image(stdMap[0,:,:] > thresh)  
    plt.title("stddev threshold = %.3f"%thresh)
```



stddev threshold = 0.007



stddev threshold = 0.010



stddev threshold = 0.015



stddev threshold = 0.020



We try several different thresholds for the standard deviation. Based on the filtering result, we choose threshold = 0.01, in which case the noise can be filtered to some extent, while the whole image can be stored at the same time.

Question 17

Filter your data such that we only keep the voxels that have the standard deviation bigger than the threshold in question 16.

```
In [42]: from numpy import std
# remove series object that has the standard deviation bigger than a threshold
filtered = data.filter(lambda x: x.std()>0.01).cache()
print(filtered.shape)
print(data.shape)

(5898, 20)
(2, 76, 87, 20)
```

Question 18


Re-train and choose the best models with different values of `K` on the new data.


```
In [43]: from sklearn.cluster import KMeans
import numpy as np
# declare the possible values of K
ks = [5, 10, 15, 20, 30, 50, 100, 200]

def buildModels(data):
    # declare the collection of models
    models = []
    # build model for each K and append to models
    for k in ks:
        models.append(KMeans(n_clusters=k).fit(data.reshape(data.count(),20)))

        #models.append(KMeans.train(data, k, maxIterations=10, initializationMode="random"))
    return models

models = buildModels(filtered)
testAndPlotTheResult(training_data, models)
```



PUT YOUR ANSWER HERE !!!


Question 19

a) Plot the centroids of the best model with a smart color selection. b) Plot the result of the clustering algorithm by a color map of voxels. c) Comment about your figures.

```
In [65]: bestModel = models[4]
newClrs = optimize_color(bestModel.cluster_centers_)
# predict the nearest cluster id for each voxel in Series
labels = data.map(lambda x : bestModel.predict(x))
# collect data to the driver
imgLabels = labels.toArray()
```

```
In [66]: brainmap = transform(newClrs, imgLabels[0,:,:])
draw_image(brainmap)
#help(draw_image)
```

```
Out[66]: <matplotlib.image.AxesImage at 0x7fb8e0460e80>
```




In the figure above we see that the neurons are clustered into different areas. With the $k = 30$, the outline between the background and the content is clear (although there is still some noise). We only see the evident clusters less than 30, which means there are some clusters with few data points.

g. Improve the visualization by adding similarity

These maps are slightly odd because pixels that did not survive our threshold still end up colored as something. A useful trick is masking pixels based on how well they match the cluster they belong to. We can compute this using the `similarity` method of `KMeansModel`.

```
In [67]: sim = data.map(lambda x: similarity(bestModel.cluster_centers_, x))  
  
imgSim = sim.toarray()  
  
# draw the mask  
draw_image(imgSim[0,:,:], cmap='gray', clim=(0,1))
```


Out[67]: <matplotlib.image.AxesImage at 0x7fb8e0320eb8>



And, it can be used as a linear mask on the colorization output

```
In [68]: brainmap = transform(newClrs, imgLabels[0,:,:], mask=imgSim[0,:,:])  
draw_image(brainmap)
```

```
Out[68]: <matplotlib.image.AxesImage at 0x7fb8e045d240>
```



Question 20

Since in the usecase we build and test the model from the same data, it can lead to overfitting problems. To avoid that, we can divide the data into training set and testing set. Note that each neuron occurs only one time in the data. So, we can not divide the data by dividing the neurons. Instead, we can divide the states of neurons into two different sets. Let's try with this approach and show the result.

Here we split the data into training and validation set, each with the number of time nodes 10. Then we implement the k means clustering on the training set to get the different models based on different k. Then we apply the models on the validation set to plot the error rate.

```
In [69]: from sklearn.cluster import KMeans
import numpy as np
# declare the possible values of K
ks = [5, 10, 15, 20, 30, 50, 100, 200]

# convert series data to rdd of values
training_data = data.tordd().map(lambda x: np.array(x[1])).cache()

def buildModels(data):
    # declare the collection of models
    models = []
    # build model for each K and append to models
    for k in ks:
        models.append(KMeans(n_clusters=k).fit(data.reshape(data.count(),20)))

        #models.append(KMeans.train(data, k, maxIterations=10,initializationMode="random"))
    return models

models = buildModels(data)
```

Split the whole data into training and testing set randomly, with the series length of 15 and 5.

```
In [82]: seed = np.arange(20)
np.random.shuffle(seed)
train_index = seed[0:10]
test_index = seed[10:20]
train_sample = data.select(lambda x : (x in train_index))
test_sample = data.select(lambda x : (x in test_index))
```

Train the model for the new training sample and choose the best model

```
In [115]: def buildModels_cv(data):
    # declare the collection of models
    models = []
    # build model for each K and append to models
    for k in ks:
        models.append(KMeans(n_clusters=k).fit(data.reshape(data.count(),10)))

        #models.append(KMeans.train(data, k, maxIterations=10,initializationMode="random"))
    return models

cv_models = buildModels_cv(train_sample)
rdd_train_sample = train_sample.tordd().map(lambda x: np.array(x[1])).cache()
```

```


In [107]: errors_1_train = np.asarray([model_error_1(rdd_train_sample, model) for model
in cv_models])
errors_2_train = np.asarray([model_error_2(rdd_train_sample, model) for model
in cv_models])
# errors of models when using testing set
errors_1_test = np.asarray([model_error_1(rdd_test_sample, model) for model in
cv_models])
errors_2_test = np.asarray([model_error_2(rdd_test_sample, model) for model in
cv_models])
# plot the errors with each value of K
plt.figure(figsize=(20,13))
line1, = plt.plot(ks, errors_1_train / errors_1_train.sum(), 'k-o',label = 'tr
ain_error_1')
line2, = plt.plot(ks, errors_2_train / errors_2_train.sum(), 'b-v',label = 'tr
ain_error_2')
line3, = plt.plot(ks, errors_1_test / errors_1_test.sum(), 'r:o',label = 'test
_error_1')
line4, = plt.plot(ks, errors_2_test /errors_2_test.sum(), 'g:v',label = 'test_
error_2')

plt.xlabel("k",fontsize=20)
plt.ylabel("error",fontsize=20)

plt.legend(handles=[line1,line2,line3,line4])

plt.show()

```



As the figure shown above, we have 4 lines: error_1 for training data`black line` error_1 for test data`red line` error_2 for training data`blue line` error_2 for test data`green line` We consider the **black&red** lines as a group, and **blue&green** lines as another group. When the two lines in the group cross, it means that this model has the same error rate in both training and test dataset, which avoids the overfitting problem. So according to this criteria, we can see the crossing points are $k = 20$ and $k = 30$ respectively. Also when k is between $[20, 30]$, the error is relatively low. So it is considered to be a good choice to set k between $[20, 30]$

Visualize the result

```
In [108]: bestModel = cv_models[4]
newClrs = optimize_color(bestModel.cluster_centers_)
# predict the nearest cluster id for each voxel in Series
labels_train = train_sample.map(lambda x : bestModel.predict(x))
labels_test= test_sample.map(lambda x : bestModel.predict(x))
# collect data to the driver
imgLabels_train = labels.toArray()
imgLabels_test = labels.toArray()

sim_train = train_sample.map(lambda x: similarity(bestModel.cluster_centers_,
x))
sim_test = test_sample.map(lambda x: similarity(bestModel.cluster_centers_,
x))


imgSim_train = sim_train.toArray()
imgSim_test = sim_test.toArray()

brainmap_train = transform(newClrs, imgLabels_train[0,:,:],
mask=imgSim_train[0,:,:])
brainmap_test = transform(newClrs, imgLabels_test[0,:,:],
mask=imgSim_test[0,:,:])
draw_image(brainmap_train)
plt.title("brain map of train set")


draw_image(brainmap_test)
plt.title("brain map of test set")
```


Out[108]: <matplotlib.text.Text at 0x7fb97b6a9080>

brain map of train set



brain map of test set



Question 21

Is using K-Means the best choice for a clustering algorithm? Comment the choice and suggest alternatives. For example, look at [Mixture Models](https://en.wikipedia.org/wiki/Mixture_model) and, if you have time, propose an alternative clustering technique.

NOTE Mixture models will be covered in the ASI course in greater detail.

```
In [117]: from sklearn.mixture import GMM
# declare the possible values of K
ks = [5, 10, 15, 20, 30, 50, 100, 200]

# convert series data to rdd of values
training_data = data.tordd().map(lambda x: np.array(x[1])).cache()

def buildGMMModels(data):
    # declare the collection of models
    models = []
    # build model for each K and append to models
    for k in ks:
        models.append(GMM(k).fit(data.reshape(data.count(),20)))
        #models.append(KMeans.train(data, k, maxIterations=10, initializationMo
de="random"))
    return models



GMMmodels = buildGMMModels(data)
```



```
In [125]: labels
```



```
Out[125]: array([ 1,  1, 126, ...,  1,  1,  1])
```



```
In [127]: for model in GMMmodels:
          data_flattened = data.reshape(data.count(),20)
          labels = model.predict(data_flattened)

          # collect data to the driver
          imgLabels = labels.reshape(2,76,87)
          # consider the voxel of the first layers
          draw_image(imgLabels[0,:,:],cmap = 'gray')
```









The result of GMM model shows a better performance with regarding to the noise filtering. From the figures we can see that more noises(isolated points) are clustered as the background, compared with the K Means model. However, the main content of the neuron map is also destroyed.

4. Summary

We studied Thunder and its important methods to work with images, such as `Image`, `Series` and how to apply them to a use case. In the use case, we used the K-Means algorithm to cluster the neurons without prior knowledge of what a good choice of K could be. Subsequently, we introduced some techniques for improving the initially obtained results, such as removing noise and considering similarity.

References

Some of the examples in this notebook are inspired from the [documentation of Thunder \(http://docs.thunder-project.org/\)](http://docs.thunder-project.org/).